

---

## ✓ COSE474-2024F: DEEP LEARNING HW1

2022320338 데이터과학과 신민서

### ✓ 0.1 Installation

```
1 !pip install d2l==1.0.3
```

```
Collecting d2l==1.0.3
  Downloading d2l-1.0.3-py3-none-any.whl.metadata (556 bytes)
Collecting jupyter==1.0.0 (from d2l==1.0.3)
  Downloading jupyter-1.0.0-py3-none-any.whl.metadata (995 bytes)
Collecting numpy==1.23.5 (from d2l==1.0.3)
  Downloading numpy-1.23.5-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (2.3 kB)
Collecting matplotlib==3.7.2 (from d2l==1.0.3)
  Downloading matplotlib-3.7.2-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (5.6 kB)
Collecting matplotlib-inline==0.1.6 (from d2l==1.0.3)
  Downloading matplotlib-inline-0.1.6-py3-none-any.whl.metadata (2.8 kB)
Collecting requests==2.31.0 (from d2l==1.0.3)
  Downloading requests-2.31.0-py3-none-any.whl.metadata (4.6 kB)
Collecting pandas==2.0.3 (from d2l==1.0.3)
  Downloading pandas-2.0.3-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (18 kB)
Collecting scipy==1.10.1 (from d2l==1.0.3)
  Downloading scipy-1.10.1-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (58 kB)
58.9/58.9 kB 4.5 MB/s eta 0:00:00
Requirement already satisfied: notebook in /usr/local/lib/python3.10/dist-packages (from jupyter==1.0.0->d2l==1.0.3)
Collecting qtconsole (from jupyter==1.0.0->d2l==1.0.3)
  Downloading qtconsole-5.6.0-py3-none-any.whl.metadata (5.0 kB)
Requirement already satisfied: jupyter-console in /usr/local/lib/python3.10/dist-packages (from jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: nbconvert in /usr/local/lib/python3.10/dist-packages (from jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: ipykernel in /usr/local/lib/python3.10/dist-packages (from jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: ipywidgets in /usr/local/lib/python3.10/dist-packages (from jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Collecting pyparsing<3.1,>=2.3.1 (from matplotlib==3.7.2->d2l==1.0.3)
  Downloading pyparsing-3.0.9-py3-none-any.whl.metadata (4.2 kB)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: traitlets in /usr/local/lib/python3.10/dist-packages (from matplotlib-inline==0.1.6->d2l==1.0.3)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas==2.0.3->d2l==1.0.3)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas==2.0.3->d2l==1.0.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests==2.31.0->d2l==1.0.3)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests==2.31.0->d2l==1.0.3)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests==2.31.0->d2l==1.0.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests==2.31.0->d2l==1.0.3)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: ipython-genutils in /usr/local/lib/python3.10/dist-packages (from ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: ipython>=5.0.0 in /usr/local/lib/python3.10/dist-packages (from ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: jupyter-client in /usr/local/lib/python3.10/dist-packages (from ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: tornado>=4.2 in /usr/local/lib/python3.10/dist-packages (from ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: widgetsnbextension>=3.6.0 in /usr/local/lib/python3.10/dist-packages (from ipywidgets->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: jupyterlab-widgets>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from ipywidgets->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: prompt-toolkit!=3.0.0,!=3.0.1,<3.1.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from jupyter-console->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: lxml in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: defusedxml in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: entrypoints>=0.2.2 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: Jinja2>=3.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: jupyter-core>=4.7 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: jupyterlab-pygments in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: mistune<2,>=0.8.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: nbclient>=0.5.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: nbformat>=5.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: pandocfilters>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: tinycss2 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: pyzmq<25,>=17 in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: argon2-cffi in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: nest-asyncio>=1.5 in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: Send2Trash>=1.8.0 in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: terminado>=0.8.3 in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: prometheus-client in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: nbclassic>=0.4.7 in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Collecting qtpy>=2.4.0 (from qtconsole->jupyter==1.0.0->d2l==1.0.3)
  Downloading QtPy-2.4.1-py3-none-any.whl.metadata (12 kB)
Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.10/dist-packages (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
Collecting jedi>=0.16 (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
  Using cached jedi-0.19.1-py2.py3-none-any.whl.metadata (22 kB)
Requirement already satisfied: decorator in /usr/local/lib/python3.10/dist-packages (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: pickleshare in /usr/local/lib/python3.10/dist-packages (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: backcall in /usr/local/lib/python3.10/dist-packages (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: pexpect>4.3 in /usr/local/lib/python3.10/dist-packages (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: platformdirs>=2.5 in /usr/local/lib/python3.10/dist-packages (from jupyter-core>=4.7->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: notebook-shim>=0.2.3 in /usr/local/lib/python3.10/dist-packages (from nbclassic>=0.4.7->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: fastjsonschema>=2.15 in /usr/local/lib/python3.10/dist-packages (from nbformat>=5.1->nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: jsonschema>=2.6 in /usr/local/lib/python3.10/dist-packages (from nbformat>=5.1->nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: wcwidth in /usr/local/lib/python3.10/dist-packages (from prompt-toolkit!=3.0.0,!=3.0.1,<3.1.0,>=2.0.0->jupyter-console->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: ptyprocess in /usr/local/lib/python3.10/dist-packages (from terminado>=0.8.3->notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: argon2-cffi-bindings in /usr/local/lib/python3.10/dist-packages (from argon2-cffi->notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4->nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->nbconvert->jupyter==1.0.0->d2l==1.0.3)
```

Requirement already satisfied: parso<0.9.0,>=0.8.3 in /usr/local/lib/python3.10/dist-packages (from jedi<0.16->ipython3.10) (0.8.3)  
Requirement already satisfied: attrs<22.2.0 in /usr/local/lib/python3.10/dist-packages (from jsonschema<2.6->nbformat) (22.2.0)  
Requirement already satisfied: jsonschema-specifications<2023.03.6 in /usr/local/lib/python3.10/dist-packages (from jsonschema<2.6->nbformat) (2023.03.6)  
Requirement already satisfied: referencing<0.28.4 in /usr/local/lib/python3.10/dist-packages (from jsonschema<2.6->nbformat) (0.28.4)  
Requirement already satisfied: rpds-py<0.7.1 in /usr/local/lib/python3.10/dist-packages (from jsonschema<2.6->nbformat) (0.7.1)  
Requirement already satisfied: jupyter-server<3,>=1.8 in /usr/local/lib/python3.10/dist-packages (from notebook-shim) (2.14.1)  
Requirement already satisfied: cffi<1.0.1 in /usr/local/lib/python3.10/dist-packages (from argon2-cffi-bindings->argon2-cffi) (1.0.1)  
Requirement already satisfied: pycparser in /usr/local/lib/python3.10/dist-packages (from cffi<1.0.1->argon2-cffi-bindings) (2.21)  
Requirement already satisfied: anyio<4,>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from jupyter-server<3,>=1.8) (3.7.1)  
Requirement already satisfied: websocket-client in /usr/local/lib/python3.10/dist-packages (from jupyter-server<3,>=1.8) (1.7.0)  
Requirement already satisfied: sniffio<1.1 in /usr/local/lib/python3.10/dist-packages (from anyio<4,>=3.1.0->jupyter-server) (1.0.0)  
Requirement already satisfied: exceptiongroup in /usr/local/lib/python3.10/dist-packages (from anyio<4,>=3.1.0->jupyter-server) (1.2.0)  
Downloading d2l-1.0.3-py3-none-any.whl (111 kB)

111.7/111.7 kB 9.1 MB/s eta 0:00:00  
Downloading jupyter-1.0.0-py2.py3-none-any.whl (2.7 kB)  
Downloading matplotlib-3.7.2-cp310-cp310-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (11.6 MB)  
11.6/11.6 MB 83.0 MB/s eta 0:00:00  
Downloading matplotlib-inline-0.1.6-py3-none-any.whl (9.4 kB)  
Downloading numpy-1.23.5-cp310-cp310-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (17.1 MB)  
17.1/17.1 MB 83.1 MB/s eta 0:00:00  
Downloading pandas-2.0.3-cp310-cp310-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (12.3 MB)  
12.3/12.3 MB 85.6 MB/s eta 0:00:00  
Downloading requests-2.31.0-py3-none-any.whl (62 kB)  
62.6/62.6 kB 5.4 MB/s eta 0:00:00  
Downloading scipy-1.10.1-cp310-cp310-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.whl (34.4 MB)  
34.4/34.4 MB 19.0 MB/s eta 0:00:00  
Downloading pyparsing-3.0.9-py3-none-any.whl (98 kB)  
98.3/98.3 kB 7.7 MB/s eta 0:00:00  
Downloading qtconsole-5.6.0-py3-none-any.whl (124 kB)  
124.7/124.7 kB 11.2 MB/s eta 0:00:00  
Downloading QtPy-2.4.1-py3-none-any.whl (93 kB)  
93.5/93.5 kB 7.7 MB/s eta 0:00:00

Using cached jedi-0.19.1-py2.py3-none-any.whl (1.6 MB)  
Installing collected packages: requests, qtpy, pyparsing, numpy, matplotlib-inline, jedi, scipy, pandas, matplotlib, qtconsole, QtPy  
Attempting uninstall: requests  
Found existing installation: requests 2.32.3  
Uninstalling requests-2.32.3:  
Successfully uninstalled requests-2.32.3  
Attempting uninstall: pyparsing  
Found existing installation: pyparsing 3.1.4  
Uninstalling pyparsing-3.1.4:  
Successfully uninstalled pyparsing-3.1.4  
Attempting uninstall: numpy  
Found existing installation: numpy 1.26.4  
Uninstalling numpy-1.26.4:  
Successfully uninstalled numpy-1.26.4  
Attempting uninstall: matplotlib-inline  
Found existing installation: matplotlib-inline 0.1.7  
Uninstalling matplotlib-inline-0.1.7:  
Successfully uninstalled matplotlib-inline-0.1.7  
Attempting uninstall: scipy  
Found existing installation: scipy 1.13.1  
Uninstalling scipy-1.13.1:  
Successfully uninstalled scipy-1.13.1  
Attempting uninstall: pandas  
Found existing installation: pandas 2.1.4  
Uninstalling pandas-2.1.4:  
Successfully uninstalled pandas-2.1.4  
Attempting uninstall: matplotlib  
Found existing installation: matplotlib 3.7.1  
Uninstalling matplotlib-3.7.1:  
Successfully uninstalled matplotlib-3.7.1

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behavior is currently the sole responsibility of the installed application. To instead require that a package can be safely reinstalled it should use a `reinstall` or `install-if-needed` command. For example, `pip install --reinstall numpy` or `pip install --install-if-needed numpy`.  
albuquerque 0.0.16 requires numpy<1.24, but you have numpy 1.23.5 which is incompatible.  
albuquerque 1.4.15 requires numpy<1.24.0, but you have numpy 1.23.5 which is incompatible.  
bigframes 1.17.0 requires numpy<1.24.0, but you have numpy 1.23.5 which is incompatible.  
chex 0.1.86 requires numpy<1.24.1, but you have numpy 1.23.5 which is incompatible.  
google-colab 1.0.0 requires pandas<2.1.4, but you have pandas 2.0.3 which is incompatible.  
google-colab 1.0.0 requires requests<2.32.3, but you have requests 2.31.0 which is incompatible.  
mizani 0.11.4 requires pandas<2.1.0, but you have pandas 2.0.3 which is incompatible.  
pandas-stubs 2.1.4.231227 requires numpy<1.26.0; python\_version < "3.13", but you have numpy 1.23.5 which is incompatible.  
plotnine 0.13.6 requires pandas<3.0.0,>=2.1.0, but you have pandas 2.0.3 which is incompatible.  
xarray 2024.9.0 requires numpy<1.24, but you have numpy 1.23.5 which is incompatible.  
xarray 2024.9.0 requires pandas<2.1, but you have pandas 2.0.3 which is incompatible.

Successfully installed d2l-1.0.3 jedi-0.19.1 jupyter-1.0.0 matplotlib-3.7.2 matplotlib-inline-0.1.6 numpy-1.23.5 pandas-2.0.3 qtconsole-5.6.0 QtPy-2.4.1 requests-2.31.0 scipy-1.10.1

**WARNING: The following packages were previously imported in this runtime:**

[matplotlib,matplotlib\_inline,mpl\_toolkits,numpy]

**You must restart the runtime in order to use newly installed versions.**

RESTART SESSION

## ✓ 2.1 Data Manipulation

### ✓ 2.1.1. Getting Started

```
1 import torch
```

```
1 x = torch.arange(12, dtype=torch.float32)
2 #arange function returns 1-diminsinal tensor of specified range
3 x #x vectors
```

```
➦ tensor([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10., 11.])
```

```
1 x.numel() #number of elements
```

```
➦ 12
```

```
1 x.shape #shape of vector
```

```
➦ torch.Size([12])
```

```
1 X = x.reshape(3, 4)
2 X #X is a 3x4 matrix
```

```
➦ tensor([[ 0.,  1.,  2.,  3.],
          [ 4.,  5.,  6.,  7.],
          [ 8.,  9., 10., 11.]])
```

```
1 X.shape #shape of matrix
```

```
➦ torch.Size([3, 4])
```

```
1 torch.zeros((2,3,4)) #zero matrix, tensor
```

```
➦ tensor([[[0., 0., 0., 0.],
           [0., 0., 0., 0.],
           [0., 0., 0., 0.]],
          [[0., 0., 0., 0.],
           [0., 0., 0., 0.],
           [0., 0., 0., 0.]])
```

```
1 torch.ones((2,3,4)) #one matrix, tensor
```

```
➦ tensor([[[1., 1., 1., 1.],
           [1., 1., 1., 1.],
           [1., 1., 1., 1.]],
          [[1., 1., 1., 1.],
           [1., 1., 1., 1.],
           [1., 1., 1., 1.]])
```

```
1 torch.randn(3,4) #randomly distributed from standard Gaussian (normal) distribution
```

```
➦ tensor([[ 1.1764,  0.0816, -0.8972,  0.5053],
          [-0.0864,  0.4782,  0.0743, -0.2758],
          [ 0.3064, -1.9109, -0.5629,  0.0577]])
```

```
1 torch.tensor([ [2,1,4,3], [1,2,3,4], [4,3,2,1] ]) #assign exact values
```

```
➦ tensor([[2, 1, 4, 3],
          [1, 2, 3, 4],
          [4, 3, 2, 1]])
```

### ✓ 2.1.2. Indexing and Slicing (Discussion)

X[start:stop] includes start, but not stop

```
1 X[-1], X[1:3] #last row, second and third row
```

```
→ (tensor([ 8.,  9., 10., 11.]),
    tensor([[ 4.,  5.,  6.,  7.],
            [ 8.,  9., 10., 11.]])
```

```
1 X[1,2] = 17
2 X #assign value
```

```
→ tensor([[ 0.,  1.,  2.,  3.],
          [ 4.,  5., 17.,  7.],
          [ 8.,  9., 10., 11.]])
```

```
1 X[:2, :] = 12 #row 1,2 and all columns
2 X
```

```
→ tensor([[12., 12., 12., 12.],
          [12., 12., 12., 12.],
          [ 8.,  9., 10., 11.]])
```

### ✓ 2.1.3. Operations

```
1 torch.exp(x)
```

```
→ tensor([162754.7969, 162754.7969, 162754.7969, 162754.7969, 162754.7969,
          162754.7969, 162754.7969, 162754.7969, 2980.9580,  8103.0840,
          22026.4648,  59874.1406])
```

```
1 x = torch.tensor([1.0, 2, 4, 8])
2 y = torch.tensor([2,2,2,2,])
3 x+y, x-y, x*y, x/y, x**y #elementwise computation
```

```
→ (tensor([ 3.,  4.,  6., 10.]),
    tensor([-1.,  0.,  2.,  6.]),
    tensor([ 2.,  4.,  8., 16.]),
    tensor([0.5000, 1.0000, 2.0000, 4.0000]),
    tensor([ 1.,  4., 16., 64.]))
```

```
1 #concatenate multiple tensors
2 X = torch.arange(12, dtype=torch.float32).reshape((3,4))
3 Y = torch.tensor([[2.0, 1, 4, 3], [1,2,3,4],[4,3,2,1]])
4 torch.cat((X,Y), dim=0), torch.cat((X,Y),dim=1)
5 #dim=0 concatenate along rows (stack the rows)
6 #dim=1 concatenate along columns
```

```
→ (tensor([[ 0.,  1.,  2.,  3.],
          [ 4.,  5.,  6.,  7.],
          [ 8.,  9., 10., 11.],
          [ 2.,  1.,  4.,  3.],
          [ 1.,  2.,  3.,  4.],
          [ 4.,  3.,  2.,  1.]]) ,
    tensor([[ 0.,  1.,  2.,  3.,  2.,  1.,  4.,  3.],
          [ 4.,  5.,  6.,  7.,  1.,  2.,  3.,  4.],
          [ 8.,  9., 10., 11.,  4.,  3.,  2.,  1.]]) )
```

```
1 X == Y #same number at [i,j] then True
```

```
→ tensor([[False,  True, False,  True],
          [False, False, False, False],
          [False, False, False, False]])
```

```
1 X.sum() #sum all elements
```

```
→ tensor(66.)
```

### 2.1.4. Broadcasting

#### ✓ 2.1.4. Discussion

- can use broadcasting when shapes differ
- the tensor with fewer dimensions is padded until dimension matches
- a: (3,1), b: (1,2),

- since a has 1 column, it is stretched to match b's columns
- a is stretch as  
[0,0],  
[1,1],  
[2,2]

```
1 a = torch.arange(3).reshape((3,1)) #3x1 shape
2 b = torch.arange(2).reshape((1,2)) #1x2 shape
3 a,b
```

```
↔ (tensor([[0],
          [1],
          [2]]),
    tensor([[0, 1]]))
```

```
1 a+b
```

```
↔ tensor([[0, 1],
          [1, 2],
          [2, 3]])
```

### ✓ 2.1.5. Saving Memory

```
1 before = id(Y)
2 Y = Y+X
3 id(Y) == before #points to different memory address
```

```
↔ False
```

- we can use in-place operation to save memory

```
1 Z = torch.zeros_like(Y) #creates tensor X with the same shape as Y
2 print('id(Z):', id(Z)) #prints address of tensor Z
3 Z[:] = X + Y #slice all the elements of Z, assign X+Y to this slice
4 print('id(Z):', id(Z)) #updates the new values directly to Z
```

```
↔ id(Z): 133390961148128
   id(Z): 133390961148128
```

```
1 before = id(X)
2 X += Y
3 id(X) == before
```

```
↔ True
```

### ✓ 2.1.6. Conversion to other Python Objects (Discussion)

Python libraries for multi-dimensional arrays (tensors)

- NumPy (ndarray)
- PyTorch (torch.Tensor)

```
1 A = X.numpy()
2 B = torch.from_numpy(A)
3 type(A), type(B)
```

```
↔ (numpy.ndarray, torch.Tensor)
```

```
1 a = torch.tensor([3.5])
2 a, a.item(), float(a), int(a) #convert to python scalar
```

```
↔ (tensor([3.5000]), 3.5, 3.5, 3)
```

## ✓ 2.2 Data Preprocessing

### ✓ 2.2.1. Reading the Dataset

```

1 import os
2
3 os.makedirs(os.path.join '..', 'data'), exist_ok=True)
4 data_file = os.path.join '..', 'data', 'house_tiny.csv')
5 with open(data_file, 'w') as f:
6     f.write(''NumRooms,RoofType,Price
7 NA,NA,127500
8 2,NA,106000
9 4,Slate,178100
10 NA,NA,140000'')

```

```

1 import pandas as pd
2
3 data = pd.read_csv(data_file)
4 print(data)

```

```

➡ NumRooms RoofType Price
0      NaN      NaN 127500
1      2.0      NaN 106000
2      4.0    Slate 178100
3      NaN      NaN 140000

```

### ✓ 2.2.2. Data Preparation (Discussion)

- supervised learning: train models to predict a target value given some input values
- NaN are missing values, which we want to handle it via **imputation** or **deletion**
- iloc: integer-location based indexing
- RoofType originally consists of Slate and NaN
- convert this into 2 columns: RoofType\_slate and RoofType\_NaN by True/false
- use pd.get\_dummies to convert categorical data into numerical data (0,1)

```

1 inputs, targets = data.iloc[:, 0:2], data.iloc[:, 2] #target selects third column
2 inputs = pd.get_dummies(inputs, dummy_na=True)
3 print(inputs)

```

```

➡ NumRooms RoofType_Slate RoofType_nan
0      NaN          False          True
1      2.0          False          True
2      4.0           True          False
3      NaN          False          True

```

```

1 inputs = inputs.fillna(inputs.mean()) #replace NaN with mean value
2 print(inputs)

```

```

➡ NumRooms RoofType_Slate RoofType_nan
0      3.0          False          True
1      2.0          False          True
2      4.0           True          False
3      3.0          False          True

```

### ✓ 2.2.3. Conversion to the Tensor Format

```

1 import torch
2 X = torch.tensor(inputs.to_numpy(dtype=float))
3 y = torch.tensor(targets.to_numpy(dtype=float))
4 X, y

```

```

➡ (tensor([[3., 0., 1.],
          [2., 0., 1.],
          [4., 1., 0.],
          [3., 0., 1.]], dtype=torch.float64),
  tensor([127500., 106000., 178100., 140000.], dtype=torch.float64))

```

## 2.2 Discussion

- We can now load datasets into tensors and manipulate them with basic mathematical operations

### ✓ 2.3 Linear Algebra

### ✓ 2.3.1. Scalars (Discussion)

$$x \in \mathbb{R}$$

$x$  is a real-values scalars

```
1 import torch
```

```
1 x = torch.tensor(3.0)
2 y = torch.tensor(2.0)
3 x+y, x*y, x/y, x**y
```

```
⇒ (tensor(5.), tensor(6.), tensor(1.5000), tensor(9.))
```

### ✓ 2.3.2. Vectors (Discussion)

- real world application: e.g. when training a model to predict the risk of a loan defaulting, each applicant can have a vector whose components correspond to quantities like income, length of employment, etc.

```
1 x = torch.arange(3)
2 x
```

```
⇒ tensor([0, 1, 2])
```

$x_2$  denotes second element of  $x$  (scalar)

```
1 x[2]
```

```
⇒ tensor(2)
```

```
1 len(x)
```

```
⇒ 3
```

```
1 x.shape
```

```
⇒ torch.Size([3])
```

### ✓ 2.3.3. Matrices

```
1 A = torch.arange(6).reshape(3,2)
2 A
```

```
⇒ tensor([[0, 1],
          [2, 3],
          [4, 5]])
```

```
1 A.T #transpose
```

```
⇒ tensor([[0, 2, 4],
          [1, 3, 5]])
```

```
1 A = torch.tensor([[1,2,3],[2,0,4],[3,4,5]])
2 A == A.T
```

```
⇒ tensor([[True, True, True],
          [True, True, True],
          [True, True, True]])
```

### ✓ 2.3.4. Tensors (Discussion)

- use tensors for nth-order arrays
- e.g. image can be represented in 3rd-order array: height, width, channel

```
1 torch.arange(24).reshape(2,3,4)
```



```

→ tensor([[[ 0,  1,  2,  3],
            [ 4,  5,  6,  7],
            [ 8,  9, 10, 11]],

          [[12, 13, 14, 15],
            [16, 17, 18, 19],
            [20, 21, 22, 23]]])

```

### ✓ 2.3.5. Basic Properties of Tensor Arithmetic

```

1 A = torch.arange(6, dtype=torch.float32).reshape(2,3)
2 B = A.clone() #copy A to B by allocating new memory
3 A, A+B

```

```

→ (tensor([[0., 1., 2.],
           [3., 4., 5.]],
          tensor([[ 0.,  2.,  4.],
                  [ 6.,  8., 10.])))

```

- Hadamard product of two matrices

```

1 A * B

```

```

→ tensor([[ 0.,  1.,  4.],
          [ 9., 16., 25.]])

```

```

1 a = 2
2 X = torch.arange(24).reshape(2,3,4)
3 a + X, (a * X).shape

```

```

→ (tensor([[[ 2,  3,  4,  5],
             [ 6,  7,  8,  9],
             [10, 11, 12, 13]],

           [[14, 15, 16, 17],
            [18, 19, 20, 21],
            [22, 23, 24, 25]]]),
    torch.Size([2, 3, 4]))

```

### ✓ 2.3.6. Reduction

```

1 x = torch.arange(3, dtype=torch.float32)
2 x, x.sum() #add all the elements in vector x

```

```

→ (tensor([0., 1., 2.]), tensor(3.))

```

```

1 A.shape, A.sum() #sum of elements of mxn matrix A

```

```

→ (torch.Size([2, 3]), tensor(15.))

```

```

1 A.shape, A.sum(axis=0).shape #sum along the rows (Axis 0)

```

```

→ (torch.Size([2, 3]), torch.Size([3]))

```

```

1 A.shape, A.sum(axis=1).shape #sum along the columns (axis 1)

```

```

→ (torch.Size([2, 3]), torch.Size([2]))

```

```

1 A.sum(axis=[0,1]) == A.sum() #same

```

```

→ tensor(True)

```

```

1 A.mean(), A.sum()/A.numel() #mean

```

```

→ (tensor(2.5000), tensor(2.5000))

```

```

1 A.mean(axis=0), A.sum(axis=0)/A.shape[0]

```

```

→ (tensor([1.5000, 2.5000, 3.5000]), tensor([1.5000, 2.5000, 3.5000]))

```

### ✓ 2.3.7. Non-Reduction Sum

```
1 sum_A = A.sum(axis=1, keepdims=True)
2 sum_A, sum_A.shape
```

```
⇒ (tensor([[ 3.],
           [12.]]),
    torch.Size([2, 1]))
```

```
1 A / sum_A #broadcasting
```

```
⇒ tensor([[0.0000, 0.3333, 0.6667],
          [0.2500, 0.3333, 0.4167]])
```

```
1 A.cumsum(axis=0) #cumulative sum of elements A along axis 0
```

```
⇒ tensor([[0., 1., 2.],
          [3., 5., 7.]])
```

### ✓ 2.3.8. Dot Products

```
1 y = torch.ones(3, dtype = torch.float32)
2 x, y, torch.dot(x,y)
```

```
⇒ (tensor([0., 1., 2.]), tensor([1., 1., 1.]), tensor(3.))
```

```
1 torch.sum(x*y) #equivalent to dot product
```

```
⇒ tensor(3.)
```

### ✓ 2.3.9. Matrix-Vector Products (Discussion)

- torch.mv(A,x): matrix-vector multiplication
- A@x is short-hand python operator for mv

```
1 A.shape, x.shape, torch.mv(A,x), A@x
```

```
⇒ (torch.Size([2, 3]), torch.Size([3]), tensor([ 5., 14.]), tensor([ 5., 14.])))
```

### ✓ 2.3.10. Matrix-Matrix Multiplication

```
1 B = torch.ones(3,4)
2 torch.mm(A,B), A@B
```

```
⇒ (tensor([[ 3.,  3.,  3.,  3.],
           [12., 12., 12., 12.]]) ,
    tensor([[ 3.,  3.,  3.,  3.],
           [12., 12., 12., 12.]]) )
```

### ✓ 2.3.11. Norms (Discussion)

- norm of a vector tells us how big it is

```
1 u = torch.tensor([3.0, -4.0])
2 torch.norm(u)
```

```
⇒ tensor(5.)
```

```
1 torch.abs(u).sum()
```

```
⇒ tensor(7.)
```

```
1 torch.norm(torch.ones((4,9)))
```

```
⇒ tensor(6.)
```

## 2.3 Discussion

- we can now use linear algebra to deal with basic mathematical objects like scalars, vectors, matrices, tensors
- Tensors can be sliced or reduced along axis0/1, via indexing or operations
- Hadamard products
- matrix-matrix products

## ✓ 2.5 Automatic Differentiation

### ✓ 2.5.1. A Simple Function (Discussion)

- goal: differentiate  $y = 2x^T x$  with respect to column vector  $x$

$$y = 2 \sum_{i=1}^n x_i^2 \qquad y = 2x^T x$$

$$x^T x = x_1^2 + x_2^2 + \cdots + x_n^2$$

```
1 x = torch.arange(4.0)
2 x
```

```
⇒ tensor([0., 1., 2., 3.])
```

```
1 x.requires_grad_(True) #need a place to store gradient of y
2 # x=torch.arange(4.0, requires_grad=True)로도 가능
3 x.grad #default value is None
```

```
1 y = 2 * torch.dot(x,x)
2 y
```

```
⇒ tensor(28., grad_fn=<MulBackward0>)
```

```
1 y.backward() #take the gradient of y with respect to x
2 x.grad
```

```
⇒ tensor([ 0.,  4.,  8., 12.])
```

```
1 x.grad == 4*x
```

```
⇒ tensor([True, True, True, True])
```

```
1 x.grad.zero_()
2 y=x.sum()
3 y.backward()
4 x.grad
```

```
⇒ tensor([1., 1., 1., 1.])
```

### ✓ 2.5.2. Backward for Non-Scalar Variables

```
1 x.grad.zero_()
2 y = x*x
3 y.backward(gradient=torch.ones(len(y))) #Faster: y.sum().backward()
4 x.grad
```

```
⇒ tensor([0., 2., 4., 6.])
```

### ✓ 2.5.3. Detaching Computation

```
1 x.grad.zero_()
2 y = x*x
3 u = y.detach()
4 z = u*x
5
6 z.sum().backward()
7 x.grad == u
```

→ tensor([True, True, True, True])

```
1 x.grad.zero_()
2 y.sum().backward()
3 x.grad == 2*x
```

→ tensor([True, True, True, True])

### ✓ 2.5.4. Gradients and Python Control Flow

```
1 def f(a):
2     b = a*2
3     while b.norm() < 1000:
4         b = b*2
5     if b.sum() > 0:
6         c = b
7     else: c = 100*b
8     return c
```

```
1 a = torch.randn(size=(), requires_grad=True)
2 d = f(a)
3 d.backward()
```

```
1 a.grad == d/a
```

→ tensor(False)

## 2.5 Discussion

- we can now compute differentiation of functions using python tools
- 1. attach gradients to those variables with respect to which we desire derivatives
- 2. record the computation of the target value
- 3. execute the backpropagation function
- 4. access the resulting gradient
- this can come handy when optimizing models

### ✓ 3.1 Linear Regression

#### ✓ 3.1.2. Vectorization for Speed

```
1 %matplotlib inline
2 import math
3 import time
4 import numpy as np
5 import torch
6 from d2l import torch as d2l
```

```
1 n = 10000 #10000-dimensional vectors
2 a = torch.ones(n)
3 b = torch.ones(n)
```

```
1 c = torch.zeros(n)
2 t = time.time()
```

```

3 for i in range(n):
4     c[i] = a[i] + b[i]
5 f'{time.time() - t:.5f} sec'

```

→ '0.13676 sec'

```

1 t = time.time()
2 d = a + b
3 f'{time.time() - t:.5f} sec' #faster method

```

→ '0.00130 sec'


### ✓ 3.1.3. The Normal distribution and squared loss

```

1 def normal(x, mu, sigma):
2     p = 1 / math.sqrt(2*math.pi*sigma**2)
3     return p*np.exp(-0.5 * (x-mu)**2 / sigma**2)

1 x = np.arange(-7,7,0.01)
2
3 params = [(0,1), (0,2), (3,1)]
4 d2l.plot(x, [normal(x, mu, sigma) for mu, sigma in params],
5           xlabel='x', ylabel='p(x)', figsize=(4.5, 2.5),
6           legend=[f'mean{mu}, std{sigma}' for mu, sigma in params])

```



- Note that changing the mean corresponds to a shift along the axis, and increasing the variance spreads the distribution out, lowering its peak.

## 3.1 Discussion

- we now can interpret traditional linear regression, and how linear models could be represented by simple neural networks

## ✓ 3.2 Object-Oriented Design for Implementation

```

1 import time
2 import numpy as np
3 import torch
4 from torch import nn
5 from d2l import torch as d2l

```

### ✓ 3.2.1. Utilities

```

1 def add_to_class(Class):
2     def wrapper(obj):
3         setattr(Class, obj.__name__, obj)
4     return wrapper

```

```

1 class A:
2     def __init__(self):
3         self.b = 1
4
5 a = A()

```

```

1 @add_to_class(A)
2 def do(self):
3     print('Class attribute "b" is', self.b)
4
5 a.do()

```

↗ Class attribute "b" is 1

```

1 class HyperParameters:
2     def save_hyperparameters(self, ignore=[]):
3         raise NotImplementedError

```

```

1 class B(d2l.HyperParameters):
2     def __init__(self, a, b, c):
3         self.save_hyperparameters(ignore=['c'])
4         print('self.a =', self.a, 'type:', type(self.a))
5         print('There is no self.c = ', not hasattr(self, 'c'))
6
7 b = B(a=1, b=2, c=3)

```

↗ self.a = 1 type: <class 'int'>  
There is no self.c = True

```


1 class ProgressBoard(d2l.HyperParameters):
2     def __init__(self, xlabel=None, ylabel=None, xlim=None,
3                 ylim=None, xscale='linear', yscale='linear',
4                 ls=['-', '--', '-.', ':'], colors=['C0', 'C1', 'C2', 'C3'],
5                 fig=None, axes=None, figsize=(3.5, 2.5), display=True):
6         self.save_hyperparameters()
7     def draw(self, x, y, label, every_n=1):
8         raise NotImplementedError

```

```

1 board = d2l.ProgressBoard('x')
2 for x in np.arange(0, 10, 0.1):
3     board.draw(x, np.sin(x), 'sin', every_n=2)
4     board.draw(x, np.cos(x), 'cos', every_n=10)

```



### ✓ 3.2.2. Models

```

1 class Module(nn.Module, d2l.HyperParameters):
2     """The base class of models."""
3     def __init__(self, plot_train_per_epoch=2, plot_valid_per_epoch=1):
4         super().__init__()
5         self.save_hyperparameters()
6         self.board = ProgressBoard()
7
8     def loss(self, y_hat, y):
9         raise NotImplementedError
10
11     def forward(self, X):
12         assert hasattr(self, 'net'), 'Neural network is defined'
13         return self.net(X)
14
15     def plot(self, key, value, train):
16         """Plot a point in animation."""
17         assert hasattr(self, 'trainer'), 'Trainer is not initied'
18         self.board.xlabel = 'epoch'
19         if train:
20             x = self.trainer.train_batch_idx / \
21                 self.trainer.num_train_batches
22             n = self.trainer.num_train_batches / \

```

```

23         self.plot_train_per_epoch
24     else:
25         x = self.trainer.epoch + 1
26         n = self.trainer.num_val_batches / \
27             self.plot_valid_per_epoch
28         self.board.draw(x, value.to(d2l.cpu()).detach().numpy(),
29                         ('train_' if train else 'val_') + key,
30                         every_n=int(n))
31
32     def training_step(self, batch):
33         l = self.loss(self(*batch[:-1]), batch[-1])
34         self.plot('loss', l, train=True)
35         return l
36
37     def validation_step(self, batch):
38         l = self.loss(self(*batch[:-1]), batch[-1])
39         self.plot('loss', l, train=False)
40
41     def configure_optimizers(self):
42         raise NotImplementedError

```

### ✓ 3.2.3. Data

```

1 class DataModule(d2l.HyperParameters):
2     """The base class of data."""
3     def __init__(self, root='./data', num_workers=4):
4         self.save_hyperparameters()
5
6     def get_dataloader(self, train):
7         raise NotImplementedError
8
9     def train_dataloader(self):
10         return self.get_dataloader(train=True)
11
12     def val_dataloader(self):
13         return self.get_dataloader(train=False)

```

### ✓ 3.2.4. Training

```

1 class Trainer(d2l.HyperParameters):
2     """The base class for training models with data."""
3     def __init__(self, max_epochs, num_gpus=0, gradient_clip_val=0):
4         self.save_hyperparameters()
5         assert num_gpus == 0, 'No GPU support yet'
6
7     def prepare_data(self, data):
8         self.train_dataloader = data.train_dataloader()
9         self.val_dataloader = data.val_dataloader()
10        self.num_train_batches = len(self.train_dataloader)
11        self.num_val_batches = (len(self.val_dataloader)
12                                if self.val_dataloader is not None else 0)
13
14    def prepare_model(self, model):
15        model.trainer = self
16        model.board.xlim = [0, self.max_epochs]
17        self.model = model
18
19    def fit(self, model, data):
20        self.prepare_data(data)
21        self.prepare_model(model)
22        self.optim = model.configure_optimizers()
23        self.epoch = 0
24        self.train_batch_idx = 0
25        self.val_batch_idx = 0
26        for self.epoch in range(self.max_epochs):
27            self.fit_epoch()
28
29    def fit_epoch(self):
30        raise NotImplementedError

```

## 4.1 Softmax Regression (Discussion)

- optimizing over discrete output spaces using probabilistic approach and treating the categories as instances of draws from a probability distribution

## ✓ 4.2 The Image Classification Dataset

```
1 %matplotlib inline
2 import time
3 import torch
4 import torchvision
5 from torchvision import transforms
6 from d2l import torch as d2l
7
8 d2l.use_svg_display()
```

### ✓ 4.2.1. Loading the Dataset

```
1 class FashionMNIST(d2l.DataModule):
2     """The Fashion-MNIST dataset."""
3     def __init__(self, batch_size=64, resize=(28, 28)):
4         super().__init__()
5         self.save_hyperparameters()
6         trans = transforms.Compose([transforms.Resize(resize),
7                                     transforms.ToTensor()])
8         self.train = torchvision.datasets.FashionMNIST(
9             root=self.root, train=True, transform=trans, download=True)
10        self.val = torchvision.datasets.FashionMNIST(
11            root=self.root, train=False, transform=trans, download=True)
```

```
1 data = FashionMNIST(resize=(32, 32))
2 len(data.train), len(data.val)
```

```
➦ Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz to ../data/FashionM
100%|██████████| 26421880/26421880 [00:07<00:00, 3447323.46it/s]
Extracting ../data/FashionMNIST/raw/train-images-idx3-ubyte.gz to ../data/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz to ../data/FashionM
100%|██████████| 29515/29515 [00:00<00:00, 324950.80it/s]
Extracting ../data/FashionMNIST/raw/train-labels-idx1-ubyte.gz to ../data/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz to ../data/FashionM
100%|██████████| 4422102/4422102 [00:01<00:00, 2909657.39it/s]
Extracting ../data/FashionMNIST/raw/t10k-images-idx3-ubyte.gz to ../data/FashionMNIST/raw

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz to ../data/FashionM
100%|██████████| 5148/5148 [00:00<00:00, 4871903.65it/s]Extracting ../data/FashionMNIST/raw/t10k-labels-idx1-ubyte.g
```

```
(60000, 10000)
```

```
1 data.train[0][0].shape
```

```
➦ torch.Size([1, 32, 32])
```

```
1 torch.Size([1, 32, 32])
```

```
➦ torch.Size([1, 32, 32])
```

```
1 @d2l.add_to_class(FashionMNIST)
2 def text_labels(self, indices):
3     """Return text labels."""
4     labels = ['t-shirt', 'trouser', 'pullover', 'dress', 'coat',
5              'sandal', 'shirt', 'sneaker', 'bag', 'ankle boot']
6     return [labels[int(i)] for i in indices]
```

### ✓ 4.2.2. Reading a Minibatch

```
1 @d2l.add_to_class(FashionMNIST)
2 def get_dataloader(self, train):
```



```

3 data = self.train if train else self.val
4 return torch.utils.data.DataLoader(data, self.batch_size, shuffle=train,
5                                     num_workers=self.num_workers)

```

```

1 X, y = next(iter(data.train_dataloader()))
2 print(X.shape, X.dtype, y.shape, y.dtype)

```

```

→ /usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create
  warnings.warn(_create_warning_msg(
    torch.Size([64, 1, 32, 32]) torch.float32 torch.Size([64]) torch.int64

```

```

1 tic = time.time()
2 for X, y in data.train_dataloader():
3     continue
4 f'{time.time() - tic:.2f} sec'

```

```

→ '15.23 sec'

```

### 4.2.3. Visualization

```


1 def show_images(imgs, num_rows, num_cols, titles=None, scale=1.5):
2     """Plot a list of images."""
3     raise NotImplementedError

```

```

1 @d2l.add_to_class(FashionMNIST)
2 def visualize(self, batch, nrows=1, ncols=8, labels=[]):
3     X, y = batch
4     if not labels:
5         labels = self.text_labels(y)
6     d2l.show_images(X.squeeze(1), nrows, ncols, titles=labels)
7 batch = next(iter(data.val_dataloader()))
8 data.visualize(batch)

```



## 4.2 Discussion

- using the Fashion-MNIST dataset, we have classified into 10 categories
- there is just one channel as the images are grayscale

## 4.3 The Base Classification Model

### 4.3.1. The classifier class

```

1 class Classifier(d2l.Module):
2     """The base class of classification models."""
3     def validation_step(self, batch):
4         Y_hat = self(*batch[:-1])
5         self.plot('loss', self.loss(Y_hat, batch[-1]), train=False)
6         self.plot('acc', self.accuracy(Y_hat, batch[-1]), train=False)

```

```

1 @d2l.add_to_class(d2l.Module)
2 def configure_optimizers(self):
3     return torch.optim.SGD(self.parameters(), lr=self.lr)

```

### 4.3.2. Accuracy

```

1 @d2l.add_to_class(Classifier)
2 def accuracy(self, Y_hat, Y, averaged=True):
3     """Compute the number of correct predictions."""
4     Y_hat = Y_hat.reshape((-1, Y_hat.shape[-1]))

```

```

5 preds = Y_hat.argmax(axis=1).type(Y.dtype)
6 compare = (preds == Y.reshape(-1)).type(torch.float32)
7 return compare.mean() if averaged else compare

```

## ✓ 4.4 Softmax Regression Implementation from Scratch

### ✓ 4.4.1 .The Softmax

```

1 X = torch.tensor([[1.0, 2.0, 3.0], [4.0, 5.0, 6.0]])
2 X.sum(0, keepdims=True), X.sum(1, keepdims=True)

```

```

→ (tensor([[5., 7., 9.]],
    tensor([[ 6.],
            [15.]])

```

```

1 def softmax(X):
2     X_exp = torch.exp(X)
3     partition = X_exp.sum(1, keepdims=True)
4     return X_exp / partition # The broadcasting mechanism is applied here

```

```

1 X = torch.rand((2, 5))
2 X_prob = softmax(X)
3 X_prob, X_prob.sum(1)

```

```

→ (tensor([[0.1647, 0.1440, 0.1487, 0.3321, 0.2106],
           [0.1333, 0.1750, 0.2688, 0.1655, 0.2574]]),
    tensor([1.0000, 1.0000]))

```

### ✓ 4.4.2. The Model

```

1 class SoftmaxRegressionScratch(d2l.Classifier):
2     def __init__(self, num_inputs, num_outputs, lr, sigma=0.01):
3         super().__init__()
4         self.save_hyperparameters()
5         self.W = torch.normal(0, sigma, size=(num_inputs, num_outputs),
6                                           requires_grad=True)
7         self.b = torch.zeros(num_outputs, requires_grad=True)
8
9     def parameters(self):
10        return [self.W, self.b]

```

```

1 @d2l.add_to_class(SoftmaxRegressionScratch)
2 def forward(self, X):
3     X = X.reshape((-1, self.W.shape[0]))
4     return softmax(torch.matmul(X, self.W) + self.b)

```

### ✓ 4.4.3. The Cross-Entropy Loss

```

1 y = torch.tensor([0, 2])
2 y_hat = torch.tensor([[0.1, 0.3, 0.6], [0.3, 0.2, 0.5]])
3 y_hat[[0, 1], y]

```

```

→ tensor([0.1000, 0.5000])

```

```

1 def cross_entropy(y_hat, y):
2     return -torch.log(y_hat[list(range(len(y_hat))), y]).mean()
3
4 cross_entropy(y_hat, y)

```

```

→ tensor(1.4979)

```

```

1 @d2l.add_to_class(SoftmaxRegressionScratch)
2 def loss(self, y_hat, y):
3     return cross_entropy(y_hat, y)


```

### ✓ 4.4.4. training

```

1 data = d2l.FashionMNIST(batch_size=256)
2 model = SoftmaxRegressionScratch(num_inputs=784, num_outputs=10, lr=0.1)
3 trainer = d2l.Trainer(max_epochs=10)
4 trainer.fit(model, data)

```



#### 4.4.5. Prediction

```

1 X, y = next(iter(data.val_data_loader()))
2 preds = model(X).argmax(axis=1)
3 preds.shape


```

```
torch.Size([256])
```

```

1 wrong = preds.type(y.dtype) != y
2 X, y, preds = X[wrong], y[wrong], preds[wrong]
3 labels = [a+'\n'+b for a, b in zip(
4     data.text_labels(y), data.text_labels(preds))]
5 data.visualize([X, y], labels=labels)

```



#### 4.4 Discussion

- we can adjust the hyperparameters such as epochs, minibatch size, and learning rate to train the model

#### 5.1 Multilayer Perceptrons

```

1 %matplotlib inline
2 import torch
3 from d2l import torch as d2l

```

##### 5.1.1.4. Universal Approximators (discussion)

- Even with single-hidden-layer network, given enough nodes (possibly absurdly many), and the right set of weights, we can model any function
- kernel methods are way more effective, since they are capable of solving the problem exactly even in **infinite dimensional spaces** (=universal approximator theorem)
- we can approximate many functions much more compactly by using deeper (rather than wider) networks


##### 5.1.2.1. ReLU Function (Discussion)

$$\text{ReLU}(x) = \max(x, 0)$$

```

1 x = torch.arange(-8.0, 8.0, 0.1, requires_grad=True)
2 y = torch.relu(x) #output of relu
3 d2l.plot(x.detach(), y.detach(), 'x', 'relu(x)', figsize=(5,2.5))


```



```

1 y.backward(torch.ones_like(x), retain_graph=True)
2 d2l.plot(x.detach(), x.grad, 'x', 'grad of relu', figsize=(5, 2.5))

```



#### 5.1.2.1. Exercise / discussion

$$\text{pReLU}(x) = \max(0, x) + \alpha \min(0, x)$$

```


1 pReLU = lambda x, a: torch.max(torch.tensor(0), x) + a * torch.min(torch.tensor(0), x)

```

```

1 y = pReLU(x=x, a=0.1)
2
3 d2l.plot(x.detach(), y.detach(), 'x', 'pReLU', figsize=(5, 2.5))


```



```


1 x.grad.data.zero_() #
2 y.backward(torch.ones_like(x), retain_graph=True)
3 d2l.plot(x.detach(), x.grad, 'x', 'grad of pReLU', figsize=(5, 2.5))

```




### 5.1.2.2. Sigmoid Function

```
1 y = torch.sigmoid(x)
2 d2l.plot(x.detach(), y.detach(), 'x', 'sigmoid(x)', figsize=(5, 2.5))
```



```
1 x.grad.data.zero_()
2 y.backward(torch.ones_like(x), retain_graph=True)
3 d2l.plot(x.detach(), x.grad, 'x', 'grad of sigmoid', figsize=(5, 2.5))
```




- Gradient Vanishing Problem ~ Backpropagation / Activation functions

### 5.1.2.3. Tanh Function


#### 5.1.2.3. Tanh Discussion

- $\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)} = \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)}$

```
1 y = torch.tanh(x)
2 d2l.plot(x.detach(), y.detach(), 'x', 'tanh(x)', figsize=(5, 2.5))
```



```
1 # Clear out previous gradients
2 x.grad.data.zero_()
3 y.backward(torch.ones_like(x), retain_graph=True)
4 d2l.plot(x.detach(), x.grad, 'x', 'grad of tanh', figsize=(5, 2.5))
```



## 5.1 Discussion

- we can incorporate nonlinear models to build more multilayer neural network architectures

## ✓ 5.2 Implementation of Multilayer Perceptrons

```
1 import torch
2 from torch import nn
3 from d2l import torch as d2l
```

### 5.2.1. Implementation from Scratch

#### ✓ 5.2.1.1. Initializing Model Parameters (Discussion)

- Fashion-MNIST has 10 classes, each image with  $28 \times 28 = 784$  grid of grayscale pixel values
- This means: 784 input features & 10 classes
- Implement 1 hidden layer and 256 hidden units

```
1 class MLPScratch(d2l.Classifier):
2     def __init__(self, num_inputs, num_outputs, num_hiddens, lr, sigma=0.01):
3         super().__init__()
4         self.save_hyperparameters()
5         self.W1 = nn.Parameter(torch.randn(num_inputs, num_hiddens) * sigma)
6         self.b1 = nn.Parameter(torch.zeros(num_hiddens))
7         self.W2 = nn.Parameter(torch.randn(num_hiddens, num_outputs) * sigma)
8         self.b2 = nn.Parameter(torch.zeros(num_outputs))
```

#### ✓ 5.2.1.2. Model (Discussion)

- Compared to softmax regression, this has 2 fully connected layers (one hidden layer, one output layer)

```
1 def relu(X):
2     a = torch.zeros_like(X)
3     return torch.max(X, a)
```

```
1 @d2l.add_to_class(MLPScratch)
2 def forward(self, X):
3     X = X.reshape((-1, self.num_inputs))
4     H = relu(torch.matmul(X, self.W1) + self.b1)
5     return torch.matmul(H, self.W2) + self.b2
```

#### ✓ 5.2.1.3. Training