

# COSE474-2024F: DEEP LEARNING HW2

2022320338 데이터과학과 신민서

## ✓ Chp7. Convolutional Neural Network

### ✓ 7.1. From Fully Connected Layers to Convolutions

#### 7.1.1 Discussion

- CNNs systemize the idea of spatial invariance, learning useful representations with fewer paramters

Overall design of the network:

- Translation invariance: network should respond similarly to the same patch, no matter where in the image
- Locality: earliest layers should focus on local regions -> aggregate to make predictions at the whole image level
- deeper layers should be able to capture longer-range features

#### 7.1.2. Discussion

$$\begin{aligned} [\mathbf{H}]_{i,j} &= [\mathbf{U}]_{i,j} + \sum_k \sum_l [\mathbf{W}]_{i,j,k,l} [\mathbf{X}]_{k,l} \\ &= [\mathbf{U}]_{i,j} + \sum_a \sum_b [\mathbf{V}]_{i,j,a,b} [\mathbf{X}]_{i+a,j+b} . \end{aligned}$$

- Translation Invariance:

$$[\mathbf{H}]_{i,j} = u + \sum_a \sum_b [\mathbf{V}]_{a,b} [\mathbf{X}]_{i+a,j+b}$$

- Locality:

$$[\mathbf{H}]_{i,j} = u + \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} [\mathbf{V}]_{a,b} [\mathbf{X}]_{i+a,j+b}$$

#### 7.1.3. Discussion

$$(f * g)(i) = \sum_a f(a)g(i - a)$$

#### 7.1.4. Discussion

- images are third-order tensors: H x W x channel (RGB):

$$[\mathbf{V}]_{a,b,c}$$

- To support multiple channels in both inputs (X) and hidden representations (H), add a fourth coordinate:

$$[\mathbf{V}]_{a,b,c,d}$$

- This makes:

$$[\mathbf{H}]_{i,j,d} = u + \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} \sum_c [\mathbf{V}]_{a,b,c,d} [\mathbf{X}]_{i+a,j+b}$$

## ✓ 7.2. Convolutions for Images

```
1 !pip install d2l==1.0.3
```

```
Collecting d2l==1.0.3
  Downloading d2l-1.0.3-py3-none-any.whl.metadata (556 bytes)
Collecting jupyter==1.0.0 (from d2l==1.0.3)
  Downloading jupyter-1.0.0-py3-none-any.whl.metadata (995 bytes)
Collecting numpy==1.23.5 (from d2l==1.0.3)
  Downloading numpy-1.23.5-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (2.3 kB)
Collecting matplotlib==3.7.2 (from d2l==1.0.3)
  Downloading matplotlib-3.7.2-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (5.6 kB)
Collecting matplotlib-inline==0.1.6 (from d2l==1.0.3)
  Downloading matplotlib-inline-0.1.6-py3-none-any.whl.metadata (2.8 kB)
Collecting requests==2.31.0 (from d2l==1.0.3)
  Downloading requests-2.31.0-py3-none-any.whl.metadata (4.6 kB)
Collecting pandas==2.0.3 (from d2l==1.0.3)
  Downloading pandas-2.0.3-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (18 kB)
Collecting scipy==1.10.1 (from d2l==1.0.3)
  Downloading scipy-1.10.1-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (58 kB)
58.9/58.9 kB 2.6 MB/s eta 0:00:00
Requirement already satisfied: notebook in /usr/local/lib/python3.10/dist-packages (from jupyter==1.0.0->d2l==1.0.3)
Collecting qtconsole (from jupyter==1.0.0->d2l==1.0.3)
  Downloading qtconsole-5.6.0-py3-none-any.whl.metadata (5.0 kB)
Requirement already satisfied: jupyter-console in /usr/local/lib/python3.10/dist-packages (from jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: nbconvert in /usr/local/lib/python3.10/dist-packages (from jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: ipykernel in /usr/local/lib/python3.10/dist-packages (from jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: ipywidgets in /usr/local/lib/python3.10/dist-packages (from jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Collecting pyparsing<3.1,>=2.3.1 (from matplotlib==3.7.2->d2l==1.0.3)
  Downloading pyparsing-3.0.9-py3-none-any.whl.metadata (4.2 kB)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: traitlets in /usr/local/lib/python3.10/dist-packages (from matplotlib-inline==0.1.6->d2l==1.0.3)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas==2.0.3->d2l==1.0.3)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas==2.0.3->d2l==1.0.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests==2.31.0->d2l==1.0.3)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests==2.31.0->d2l==1.0.3)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests==2.31.0->d2l==1.0.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests==2.31.0->d2l==1.0.3)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib==3.7.2->d2l==1.0.3)
Requirement already satisfied: ipython-genutils in /usr/local/lib/python3.10/dist-packages (from ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: ipython>=5.0.0 in /usr/local/lib/python3.10/dist-packages (from ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: jupyter-client in /usr/local/lib/python3.10/dist-packages (from ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: tornado>=4.2 in /usr/local/lib/python3.10/dist-packages (from ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: widgetsnbextension>=3.6.0 in /usr/local/lib/python3.10/dist-packages (from ipywidgets->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: jupyterlab-widgets>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from ipywidgets->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: prompt-toolkit!=3.0.0,!<3.1.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from jupyter-console->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: pygments in /usr/local/lib/python3.10/dist-packages (from jupyter-console->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: lxml in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: defusedxml in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: entrypoints>=0.2.2 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: Jinja2>=3.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: jupyter-core>=4.7 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: jupyterlab-pygments in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: mistune<2,>=0.8.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: nbclient>=0.5.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: nbformat>=5.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: pandocfilters>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: tinycss2 in /usr/local/lib/python3.10/dist-packages (from nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: pyzmq<25,>=17 in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: argon2-cffi in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: nest-asyncio>=1.5 in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: Send2Trash>=1.8.0 in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: terminado>=0.8.3 in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: prometheus-client in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: nbclassic>=0.4.7 in /usr/local/lib/python3.10/dist-packages (from notebook->jupyter==1.0.0->d2l==1.0.3)
Collecting qtpy>=2.4.0 (from qtconsole->jupyter==1.0.0->d2l==1.0.3)
  Downloading QtPy-2.4.1-py3-none-any.whl.metadata (12 kB)
Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.10/dist-packages (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
Collecting jedi>=0.16 (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
  Using cached jedi-0.19.1-py2.py3-none-any.whl.metadata (22 kB)
Requirement already satisfied: decorator in /usr/local/lib/python3.10/dist-packages (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: pickleshare in /usr/local/lib/python3.10/dist-packages (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: backcall in /usr/local/lib/python3.10/dist-packages (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: pexpect>4.3 in /usr/local/lib/python3.10/dist-packages (from ipython>=5.0.0->ipykernel->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: platformdirs>=2.5 in /usr/local/lib/python3.10/dist-packages (from jupyter-core>=4.7->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: notebook-shim>=0.2.3 in /usr/local/lib/python3.10/dist-packages (from nbclassic>=0.4.7->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: fastjsonschema>=2.15 in /usr/local/lib/python3.10/dist-packages (from nbformat>=5.1->nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: jsonschema>=2.6 in /usr/local/lib/python3.10/dist-packages (from nbformat>=5.1->nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: wcwidth in /usr/local/lib/python3.10/dist-packages (from prompt-toolkit!=3.0.0,!<3.1.0,>=2.0.0->jupyter-console->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: ptyprocess in /usr/local/lib/python3.10/dist-packages (from terminado>=0.8.3->notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: argon2-cffi-bindings in /usr/local/lib/python3.10/dist-packages (from argon2-cffi->notebook->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: soupsieve>=1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4->nbconvert->jupyter==1.0.0->d2l==1.0.3)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from bleach->nbconvert->jupyter==1.0.0->d2l==1.0.3)
```

```
Requirement already satisfied: parso<0.9.0,>=0.8.3 in /usr/local/lib/python3.10/dist-packages (from jedi<0.19.1,py>=3.10,py3-none-any.whl)
Requirement already satisfied: attrs>=22.2.0 in /usr/local/lib/python3.10/dist-packages (from jsonschema>=2.6->nbformat)
Requirement already satisfied: jsonschema-specifications>=2023.03.6 in /usr/local/lib/python3.10/dist-packages (from jsonschema>=2.6->nbformat)
Requirement already satisfied: referencing>=0.28.4 in /usr/local/lib/python3.10/dist-packages (from jsonschema>=2.6->nbformat)
Requirement already satisfied: rpds-py>=0.7.1 in /usr/local/lib/python3.10/dist-packages (from jsonschema>=2.6->nbformat)
Requirement already satisfied: jupyter-server<3,>=1.8 in /usr/local/lib/python3.10/dist-packages (from notebook-shim)
Requirement already satisfied: cffi>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from argon2-cffi-bindings->argon2-cffi)
Requirement already satisfied: pycparser in /usr/local/lib/python3.10/dist-packages (from cffi>=1.0.1->argon2-cffi-bindings)
Requirement already satisfied: anyio<4,>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from jupyter-server<3,>=1.8)
Requirement already satisfied: websocket-client in /usr/local/lib/python3.10/dist-packages (from jupyter-server<3,>=1.8)
Requirement already satisfied: sniffio>=1.1 in /usr/local/lib/python3.10/dist-packages (from anyio<4,>=3.1.0->jupyter-server)
Requirement already satisfied: exceptiongroup in /usr/local/lib/python3.10/dist-packages (from anyio<4,>=3.1.0->jupyter-server)
Downloading d2l-1.0.3-py3-none-any.whl (111 kB)
```

```
111.7/111.7 kB 4.8 MB/s eta 0:00:00
Downloading jupyter-1.0.0-py2.py3-none-any.whl (2.7 kB)
Downloading matplotlib-3.7.2-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (11.6 MB)
11.6/11.6 MB 70.9 MB/s eta 0:00:00
Downloading matplotlib-inline-0.1.6-py3-none-any.whl (9.4 kB)
Downloading numpy-1.23.5-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (17.1 MB)
17.1/17.1 MB 66.6 MB/s eta 0:00:00
Downloading pandas-2.0.3-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (12.3 MB)
12.3/12.3 MB 78.6 MB/s eta 0:00:00
Downloading requests-2.31.0-py3-none-any.whl (62 kB)
62.6/62.6 kB 4.3 MB/s eta 0:00:00
Downloading scipy-1.10.1-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (34.4 MB)
34.4/34.4 MB 13.4 MB/s eta 0:00:00
Downloading pyparsing-3.0.9-py3-none-any.whl (98 kB)
98.3/98.3 kB 5.6 MB/s eta 0:00:00
Downloading qtconsole-5.6.0-py3-none-any.whl (124 kB)
124.7/124.7 kB 9.0 MB/s eta 0:00:00
Downloading QtPy-2.4.1-py3-none-any.whl (93 kB)
93.5/93.5 kB 5.8 MB/s eta 0:00:00
```

```
Using cached jedi-0.19.1-py2.py3-none-any.whl (1.6 MB)
Installing collected packages: requests, qtpy, pyparsing, numpy, matplotlib-inline, jedi, scipy, pandas, matplotlib,
  Attempting uninstall: requests
    Found existing installation: requests 2.32.3
    Uninstalling requests-2.32.3:
      Successfully uninstalled requests-2.32.3
  Attempting uninstall: pyparsing
    Found existing installation: pyparsing 3.1.4
    Uninstalling pyparsing-3.1.4:
      Successfully uninstalled pyparsing-3.1.4
  Attempting uninstall: numpy
    Found existing installation: numpy 1.26.4
    Uninstalling numpy-1.26.4:
      Successfully uninstalled numpy-1.26.4
  Attempting uninstall: matplotlib-inline
    Found existing installation: matplotlib-inline 0.1.7
    Uninstalling matplotlib-inline-0.1.7:
      Successfully uninstalled matplotlib-inline-0.1.7
  Attempting uninstall: scipy
    Found existing installation: scipy 1.13.1
    Uninstalling scipy-1.13.1:
      Successfully uninstalled scipy-1.13.1
  Attempting uninstall: pandas
    Found existing installation: pandas 2.2.2
    Uninstalling pandas-2.2.2:
      Successfully uninstalled pandas-2.2.2
  Attempting uninstall: matplotlib
    Found existing installation: matplotlib 3.7.1
    Uninstalling matplotlib-3.7.1:
      Successfully uninstalled matplotlib-3.7.1
```

```
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This beha
albuscore 0.0.16 requires numpy>=1.24, but you have numpy 1.23.5 which is incompatible.
alumentations 1.4.15 requires numpy>=1.24.4, but you have numpy 1.23.5 which is incompatible.
bigframes 1.21.0 requires numpy>=1.24.0, but you have numpy 1.23.5 which is incompatible.
chex 0.1.87 requires numpy>=1.24.1, but you have numpy 1.23.5 which is incompatible.
google-colab 1.0.0 requires pandas==2.2.2, but you have pandas 2.0.3 which is incompatible.
google-colab 1.0.0 requires requests==2.32.3, but you have requests 2.31.0 which is incompatible.
jax 0.4.33 requires numpy>=1.24, but you have numpy 1.23.5 which is incompatible.
jaxlib 0.4.33 requires numpy>=1.24, but you have numpy 1.23.5 which is incompatible.
mizani 0.11.4 requires pandas>=2.1.0, but you have pandas 2.0.3 which is incompatible.
plotnine 0.13.6 requires pandas<3.0.0,>=2.1.0, but you have pandas 2.0.3 which is incompatible.
xarray 2024.9.0 requires numpy>=1.24, but you have numpy 1.23.5 which is incompatible.
xarray 2024.9.0 requires pandas>=2.1, but you have pandas 2.0.3 which is incompatible.
Successfully installed d2l-1.0.3 jedi-0.19.1 jupyter-1.0.0 matplotlib-3.7.2 matplotlib-inline-0.1.6 numpy-1.23.5 pan
```

**WARNING: The following packages were previously imported in this runtime:**

**[matplotlib,matplotlib\_inline,mpl\_toolkits,numpy]**

**You must restart the runtime in order to use newly installed versions.**

RESTART SESSION



```

1 import torch
2 from torch import nn
3 from d2l import torch as d2l

```

### ✓ 7.2.1. Cross-Correlation Operation

```

1 def corr2d(X, K):
2     # Compute 2D cross-correlation
3     h, w = K.shape
4     Y = torch.zeros( (X.shape[0] - h + 1, X.shape[1] - w + 1) )
5     for i in range(Y.shape[0]):
6         for j in range(Y.shape[1]):
7             Y[i, j] = (X[i:i + h, j:j + w] * K).sum()
8     return Y

1 X = torch.tensor([[0.01, 1.0, 2.0], [3.0, 4.0, 5.0], [6.0, 7.0, 8.0]])
2 K = torch.tensor([[0.0, 1.0], [2.0, 3.0]])
3 corr2d(X, K)
4
⇒ tensor([[19., 25.],
          [37., 43.]])

```

### ✓ 7.2.2. Convolutional Layers

```

1 class Conv2D(nn.Module):
2     def __init__(self, kernel_size):
3         super().__init__()
4         self.weight = nn.Parameter(torch.rand(kernel_size))
5         self.bias = nn.Parameter(torch.zeros(1))
6
7     def forward(self, x):
8         return corr2d(x, self.weight) + self.bias
9

```

### ✓ 7.2.3. Object Edge Detection

```

1 X = torch.ones((6,8)) #image of 6x8 pixels
2 X[:, 2:6] = 0
3 X

```

```

⇒ tensor([[1., 1., 0., 0., 0., 0., 1., 1.],
          [1., 1., 0., 0., 0., 0., 1., 1.],
          [1., 1., 0., 0., 0., 0., 1., 1.],
          [1., 1., 0., 0., 0., 0., 1., 1.],
          [1., 1., 0., 0., 0., 0., 1., 1.],
          [1., 1., 0., 0., 0., 0., 1., 1.]])

```

```

1 #construct kernel K of 1x2
2 K = torch.tensor([[1.0, -1.0]])

```

```

1 Y = corr2d(X, K)
2 Y

```

```

⇒ tensor([[ 0.,  1.,  0.,  0.,  0., -1.,  0.],
          [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
          [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
          [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
          [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
          [ 0.,  1.,  0.,  0.,  0., -1.,  0.]])

```

```

1 corr2d(X.t(), K)

```

```

⇒ tensor([[0., 0., 0., 0., 0.],
          [0., 0., 0., 0., 0.],
          [0., 0., 0., 0., 0.],
          [0., 0., 0., 0., 0.],
          [0., 0., 0., 0., 0.],
          [0., 0., 0., 0., 0.]])

```

## ✓ 7.2.4. Learning a Kernel

```
1 conv2d = nn.LazyConv2d(1, kernel_size=(1,2), bias=False)
2
3 X = X.reshape((1,1,6,8))
4 Y = Y.reshape((1,1,6,7))
5
6 lr = 3e-2 #learning rate
7
8 for i in range(10):
9     Y_hat = conv2d(X)
10    l = (Y_hat - Y) ** 2
11    conv2d.zero_grad()
12    l.sum().backward()
13    conv2d.weight.data[:] -= lr * conv2d.weight.grad
14    if (i+1) % 2 == 0:
15        print(f'epoch {i+1}, loss {l.sum():.3f}')
```

```
↩ epoch 2, loss 7.683
    epoch 4, loss 1.290
    epoch 6, loss 0.217
    epoch 8, loss 0.036
    epoch 10, loss 0.006
```

```
1 conv2d.weight.data.reshape((1,2))
```

```
↩ tensor([[ 0.9847, -0.9865]])
```

## 7.2.5. Discussion

- since kernels are learned from data in DL, the outputs of convolutional layers remain unaffected no matter such layers perform either the strict convolution operations or the cross-correlation operations
- When the convolutional layer performs strict convolution for an input, the same output will be obtained (cross-correlation of the input and K)
- 

## 7.2.6 Discussion

- Feature map: convolutional layer output of features in spatial dimensions (WxH)
- Receptive field: elements from all the previous layers that may affect the calculation of x during the forward propagation

## ✓ 7.3. Padding and Stride

```
1 import torch
2 from torch import nn
```

### ✓ 7.3.1. Padding (discussion)

- Padding: adding extra pixels (usually zeros) around the edges of an input image
- Why? To prevent shrinking of the feature map
- If we add a total of  $P_h$  rows of padding and  $P_w$  columns of padding, the output shape will be:
$$(n_h - k_h + p_h + 1) * (n_w - k_w + p_w + 1)$$
- height and width increase by  $p_h$  and  $p_w$
- usually use convolution kernels with odd height and width values

```
1 def comp_conv2d(conv2d, X):
2     X = X.reshape((1,1) + X.shape)
3     Y = conv2d(X)
4     return Y.reshape(Y.shape[2:])
5
6 conv2d = nn.LazyConv2d(1, kernel_size=3, padding=1)
7 X = torch.rand(size=(8,8))
8 comp_conv2d(conv2d, X).shape
```

```
↩ torch.Size([8, 8])
```

```
1 conv2d = nn.LazyConv2d(1, kernel_size=(5,3), padding=(2,1))
2 comp_conv2d(conv2d, X).shape
```

```
↔ torch.Size([8, 8])
```

### ✓ 7.3.2. Stride (discussion)

- Stride: number of rows and columns traversed per slide
- generally, the output shape with  $S_h$  and  $S_w$  is:

$$\left\lfloor \frac{n_h - k_h + p_h + s_h}{s_h} \right\rfloor \times \left\lfloor \frac{n_w - k_w + p_w + s_w}{s_w} \right\rfloor$$

```
1 conv2d = nn.LazyConv2d(1, kernel_size=3, padding = 1, stride = 2)
2 comp_conv2d(conv2d, X).shape
```





```
↔ torch.Size([4, 4])
```

```
1 conv2d = nn.LazyConv2d(1, kernel_size=(3,5), padding=(0,1), stride=(3,4))
2 comp_conv2d(conv2d, X).shape
```

```
↔ torch.Size([2, 2])
```

## 7.3 Discussion

### • Padding

-  increases height and width of the output
-  in order to avoid undesirable shrinkage of the output
-  ensures that all pixels are used equally frequently
-  usually use symmetric padding on both sides

### • Stride

-  reduces the resolution of the output

## ✓ 7.4. Multiple Input and Multiple Output Channels

```
1 import torch
2 from d2l import torch as d2l
```

### ✓ 7.4.1. Multiple input channels (Discussion)

- If the number of channels for input data is  $c_i$ , then the number of input channels of the convolution kernel should also be  $c_i$
- Given  $c_i > 1$ , the shape is:  $c_i * k_h * k_w$
- 

```
1 def corr2d_multi_in(X,K):
2     return sum(d2l.corr2d(x, k) for x, k in zip(X,K))
```

```
1 X = torch.tensor([[[0.0,1.0,2.0],[3.0,4.0,5.0],[6.0,7.0,8.0]],[[1.0,2.0,3.0],[4.0,5.0,6.0],[7.0,8.0,9.0]]])
2 K = torch.tensor([[[0.0,1.0],[2.0,3.0]],[[1.0,2.0],[3.0,4.0]])]
3
4 corr2d_multi_in(X,K)
```

```
↔ tensor([[ 56.,  72.],
         [104., 120.]])
```

### ✓ 7.4.2. Multiple Output Channels (Discussion)

- In deeper channel network, each channel responds to different set of features

- Channels are optimized to be jointly useful (some direction in channel space corresponds to detecting edges)
- Shape of multiple channel output:  $c_o * c_i * k_h * k_w$

```
1 def corr2d_multi_in_out(X,K):
2     return torch.stack([corr2d_multi_in(X,k) for k in K], 0)
```

```
1 K = torch.stack((K, K+1, K+2), 0)
2 K.shape
```

```
→ torch.Size([3, 2, 2, 2])
```

```
1 corr2d_multi_in_out(X,K)
```

```
→ tensor([[[[ 56.,  72.],
              [104., 120.]],

            [[ 76., 100.],
              [148., 172.]],

            [[ 96., 128.],
              [192., 224.]]]])
```

### ✓ 7.3.1. 1x1 Convolutional Layer (discussion)

- The only computation of the 1x1 convolution happens on the channel dimension
- 

```
1 def corr2d_multi_in_out_1x1(X,K):
2     c_i, h, w = X.shape
3     c_o = K.shape[0]
4     X = X.reshape((c_i, h*w))
5     K = K.reshape((c_o, c_i))
6     Y = torch.matmul(K,X)
7     return Y.reshape((c_o, h, w))
```

```
1 X = torch.normal(0,1, (3,3,3))
2 K = torch.normal(0,1, (2,3,1,1))
3 Y1 = corr2d_multi_in_out_1x1(X,K)
4 Y2 = corr2d_multi_in_out(X,K)
5 assert float(torch.abs(Y1-Y2).sum()) < 1e-6
```

## ✓ 7.5. Pooling

### ✓ 7.5.1. Maximum & Average Pooling (Discussion)

- pooling: technique used to reduce spatial dimensions, which helps to decrease computation
- a fixed-shape window that is slid over all regions in the input according to its stride, computing a single output for each location traversed by the fixed-shape window
- but there is no kernel
- calculates either the max or avg. value of the elements in the pooling window

```
1 def pool2d(X, pool_size, mode='max'):
2     p_h, p_w = pool_size
3     Y = torch.zeros((X.shape[0] - p_h + 1, X.shape[1] - p_w + 1))
4     for i in range(Y.shape[0]):
5         for j in range(Y.shape[1]):
6             if mode == 'max':
7                 Y[i,j] = X[i: i+p_h, j: j+p_w].max()
8             elif mode == 'avg':
9                 Y[i,j] = X[i: i+p_h, j: j+p_w].mean()
10    return Y
```

```
1 X = torch.tensor([[0.0, 1.0, 2.0], [3.0, 4.0, 5.0], [6.0, 7.0, 8.0]])
2
3 pool2d(X, (2,2))
```

```
→ tensor([[4., 5.],
          [7., 8.]])
```



```
1 pool2d(X, (2,2), 'avg')
```

```
↔ tensor([[2., 3.],  
         [5., 6.]])
```

### ✓ 7.5.2. Padding and Stride (Discussion)

- We can demonstrate the use of padding and stride in pooling layers via built-in 2D max-pooling layer from DL framework

```
1 X = torch.arange(16, dtype = torch.float32).reshape((1,1,4,4))  
2 X
```

```
↔ tensor([[[[ 0.,  1.,  2.,  3.],  
            [ 4.,  5.,  6.,  7.],  
            [ 8.,  9., 10., 11.],  
            [12., 13., 14., 15.]]]])
```

```
1 pool2d = nn.MaxPool2d(3)  
2 pool2d(X)
```

```
↔ tensor([[[[10.]]]])
```

```
1 pool2d = nn.MaxPool2d(3, padding=1, stride=2)  
2 pool2d(X)
```

```
↔ tensor([[[[ 5.,  7.],  
            [13., 15.]]]])
```

```
1 pool2d = nn.MaxPool2d((2,3), stride=(2,3), padding=(0,1))  
2 pool2d(X)
```

```
↔ tensor([[[[ 5.,  7.],  
            [13., 15.]]]])
```

### ✓ 7.5.3. Multiple Channels

```
1 X = torch.cat((X, X+1), 1)  
2 X
```

```
↔ tensor([[[[ 0.,  1.,  2.,  3.],  
            [ 4.,  5.,  6.,  7.],  
            [ 8.,  9., 10., 11.],  
            [12., 13., 14., 15.]],  
          [[ 1.,  2.,  3.,  4.],  
            [ 5.,  6.,  7.,  8.],  
            [ 9., 10., 11., 12.],  
            [13., 14., 15., 16.]]]])
```

```
1 pool2d = nn.MaxPool2d(3, padding=1, stride=2)  
2 pool2d(X)
```

```
↔ tensor([[[[ 5.,  7.],  
            [13., 15.]],  
          [[ 6.,  8.],  
            [14., 16.]]]])
```

## ✓ 7.6. Convolutional Neural Networks (LeNet)

```
1 import torch  
2 from torch import nn  
3 from d2l import torch as d2l
```

### ✓ 7.6.1. LeNet

- 2 parts in LeNet: (i) a convolutional encoder consisting of 2 convolutional layers (ii) a dense block consisting of 3 fully connected layers (120, 84, 10 outputs)

- To pass output from convolutional block to the dense block, each example in the minibatch should be flattened

```

1 def init_cnn(module):
2     if type(module) == nn.Linear or type(module) == nn.Conv2d:
3         nn.init.xavier_uniform_(module.weight)
4
5 class LeNet(d2l.Classifier):
6     def __init__(self, lr=0.1, num_classes=10):
7         super().__init__()
8         self.save_hyperparameters()
9         self.net = nn.Sequential(
10             nn.LazyConv2d(6, kernel_size=5, padding=2), nn.Sigmoid(),
11             nn.AvgPool2d(kernel_size=2, stride=2),
12             nn.LazyConv2d(16, kernel_size=5), nn.Sigmoid(),
13             nn.AvgPool2d(kernel_size=2, stride=2),
14             nn.Flatten(),
15             nn.LazyLinear(120), nn.Sigmoid(),
16             nn.LazyLinear(84), nn.Sigmoid(),
17             nn.LazyLinear(num_classes))
18

```

```

1 @d2l.add_to_class(d2l.Classifier)
2 def layer_summary(self, X_shape):
3     X = torch.randn(*X_shape)
4     for layer in self.net:
5         X = layer(X)
6         print(layer.__class__.__name__, 'output shape:\t', X.shape)
7
8 model = LeNet()
9 model.layer_summary((1,1,28,28))

```

```

→ Conv2d output shape:      torch.Size([1, 6, 28, 28])
  Sigmoid output shape:     torch.Size([1, 6, 28, 28])
  AvgPool2d output shape:   torch.Size([1, 6, 14, 14])
  Conv2d output shape:      torch.Size([1, 16, 10, 10])
  Sigmoid output shape:     torch.Size([1, 16, 10, 10])
  AvgPool2d output shape:   torch.Size([1, 16, 5, 5])
  Flatten output shape:     torch.Size([1, 400])
  Linear output shape:      torch.Size([1, 120])
  Sigmoid output shape:     torch.Size([1, 120])
  Linear output shape:      torch.Size([1, 84])
  Sigmoid output shape:     torch.Size([1, 84])
  Linear output shape:      torch.Size([1, 10])


```

## 7.6.2. Training

```

1 trainer = d2l.Trainer(max_epochs = 10, num_gpus = 1)
2 data = d2l.FashionMNIST(batch_size=128)
3 model = LeNet(lr=0.1)
4 model.apply_init([next(iter(data.get_dataloader(True)))[0]), init_cnn]
5 trainer.fit(model, data)

```



## Chp8. Modern Convolutional Neural Networks

### 8.2 Networks Using Blocks (VGG)

```

1 import torch
2 from torch import nn

```

```
3 from d2l import torch as d2l
```

### ✓ 8.2.1. VGG Blocks (Discussion)

- Problem of basic building block of CNNs is that the spatial resolution decreases quite rapidly
- VGG block: consists of a sequence of convolutions with 3x3 kernels with padding of 1, and 2x2 max-pooling layer with stride of 2

```
1 def vgg_block(num_convs, out_channels):
2     layers = []
3     for _ in range(num_convs):
4         layers.append(nn.Conv2d(out_channels, kernel_size=3, padding=1))
5         layers.append(nn.ReLU())
6     layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
7     return nn.Sequential(*layers)
```

### ✓ 8.2.2. VGG Network (Discussion)

- VGG Network is divided into 2 parts: (i) consists mostly of convolutional and pooling layers (ii) consists of fully connected layers
- We need to build a specific network we simply iterate over arch to compose the blocks

```
1 class VGG(d2l.Classifier):
2     def __init__(self, arch, lr=0.1, num_classes=10):
3         super().__init__()
4         self.save_hyperparameters()
5         conv_blks = []
6         for (num_convs, out_channels) in arch:
7             conv_blks.append(vgg_block(num_convs, out_channels))
8         self.net = nn.Sequential(
9             *conv_blks, nn.Flatten(),
10             nn.Linear(4096), nn.ReLU(), nn.Dropout(0.5),
11             nn.Linear(4096), nn.ReLU(), nn.Dropout(0.5),
12             nn.Linear(num_classes))
13         self.net.apply(d2l.init_cnn)

1 VGG(arch=((1, 64), (1, 128), (2, 256), (2, 512), (2, 512))).layer_summary((1,1,224,224))
```

```
→ Sequential output shape:      torch.Size([1, 64, 112, 112])
Sequential output shape:      torch.Size([1, 128, 56, 56])
Sequential output shape:      torch.Size([1, 256, 28, 28])
Sequential output shape:      torch.Size([1, 512, 14, 14])
Sequential output shape:      torch.Size([1, 512, 7, 7])
Flatten output shape:         torch.Size([1, 25088])
Linear output shape:          torch.Size([1, 4096])
ReLU output shape:            torch.Size([1, 4096])
Dropout output shape:         torch.Size([1, 4096])
Linear output shape:          torch.Size([1, 4096])
ReLU output shape:            torch.Size([1, 4096])
Dropout output shape:         torch.Size([1, 4096])
Linear output shape:          torch.Size([1, 10])
```

### ✓ 8.2.3. Training

```
1 model = VGG(arch=((1,16), (1,32), (2, 64), (2, 128)), lr=0.01)
2 trainerr = d2l.Trainer(max_epochs=10, num_gpus=1)
3 data = d2l.FashionMNIST(batch_size=128, resize=(224,224))
4 model.apply_init([next(iter(data.get_data_loader(True)))[0]), d2l.init_cnn]
5 trainer.fit(model, data)
```

## ✓ 8.6. Residual Networks and ResNeXt

### 8.6.1 Discussion

- Say,  $F$  is the class of functions that a specific network architecture can reach
- Given a data set with features  $X$  and labels  $y$ , we can find  $f_F^*$  following optimization problem:

$$f_F^* \stackrel{\text{def}}{=} \arg \min_f L(X, y, f) \text{ subject to } f \in F$$

- Only if larger function classes contain the smaller ones -> guarantee that increasing them strictly increases the expressive power of the network

## 8.6.2 Discussion

- With residual blocks, inputs can forward propagate faster via residual connections across layers
- In the example, the residual block can be thought of as a special case of the multi-branch Inception block (2 branches one of which is the identity mapping)

```
1 import torch
2 from torch import nn
3 from torch.nn import functional as F
4 from d2l import torch as d2l

1 class Residual(nn.Module):
2     def __init__(self, num_channels, use_1x1conv=False, strides=1):
3         super().__init__()
4         self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1, stride=strides)
5         self.conv2 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1)
6         if use_1x1conv:
7             self.conv3 = nn.LazyConv2d(num_channels, kernel_size=1, stride=strides)
8         else:
9             self.conv3 = None
10        self.bn1 = nn.BatchNorm2d(num_channels)
11        self.bn2 = nn.BatchNorm2d(num_channels)
12
13    def forward(self, X):
14        Y = F.relu(self.bn1(self.conv1(X)))
15        Y = self.bn2(self.conv2(Y))
16        if self.conv3:
17            X = self.conv3(X)
18        Y += X
19        return F.relu(Y)
```

```
1 blk = Residual(3)
2 X = torch.randn(4,3,6,6)
3 blk(X).shape
```

```
↩ torch.Size([4, 3, 6, 6])
```

```
1 blk = Residual(6, use_1x1conv=True, strides=2)
2 blk(X).shape
```

```
↩ torch.Size([4, 6, 3, 3])
```

## 8.6.3 ResNet

```
1 class ResNet(d2l.Classifier):
2     def b1(self):
3         return nn.Sequential(
4             nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
5             nn.LazyBatchNorm2d(), nn.ReLU())
```