

✓ Homework 4

Instructions

- This homework focuses on understanding and applying CoCoOp for CLIP prompt tuning. It consists of **four questions** designed to assess both theoretical understanding and practical application.
- Please organize your answers and results for the questions below and submit this jupyter notebook as a **.pdf file**.
- **Deadline: 11/26 (Sat) 23:59**

➤ Preparation

- Run the code below before proceeding with the homework.
- If an error occurs, click 'Run Session Again' and then restart the runtime from the beginning.

🔍 숨겨진 셀 1개

✓ Q1. Understanding and implementing CoCoOp

- We have learned how to define CoOp in Lab Session 4.
- The main difference between CoOp and CoCoOp is **meta network** to extract image tokens that is added to the text prompt.
- Based on the CoOp code given in Lab Session 4, fill-in-the-blank exercise (4 blanks!!) to test your understanding of critical parts of the CoCoOp.

```
1 import torch.nn as nn
2
3 class CoCoOpPromptLearner(nn.Module):
4     def __init__(self, cfg, classnames, clip_model):
5         super().__init__()
6         n_cls = len(classnames)
7         n_ctx = cfg.TRAINER.COOCOOP.N_CTX
8         ctx_init = cfg.TRAINER.COOCOOP.CTX_INIT
9         dtype = clip_model.dtype
10        ctx_dim = clip_model.ln_final.weight.shape[0]
11        vis_dim = clip_model.visual.output_dim
12        clip_imgsize = clip_model.visual.input_resolution
13        cfg_imgsize = cfg.INPUT.SIZE[0]
14        assert cfg_imgsize == clip_imgsize, f"cfg_imgsize ({cfg_imgsize}) must equal to clip_imgsize ({clip_imgsize})"
15
16        if ctx_init:
17            # use given words to initialize context vectors
18            ctx_init = ctx_init.replace("_", " ")
19            n_ctx = len(ctx_init.split(" "))
20            prompt = clip.tokenize(ctx_init)
21            with torch.no_grad():
22                embedding = clip_model.token_embedding(prompt).type(dtype)
23            ctx_vectors = embedding[0, 1: 1 + n_ctx, :]
24            prompt_prefix = ctx_init
25        else:
26            # random initialization
27            ctx_vectors = torch.empty(n_ctx, ctx_dim, dtype=dtype)
28            nn.init.normal_(ctx_vectors, std=0.02)
29            prompt_prefix = " ".join(["X"] * n_ctx)
30
31        print(f'Initial context: "{prompt_prefix}")')
32        print(f"Number of context words (tokens): {n_ctx}")
33
34        self.ctx = nn.Parameter(ctx_vectors) # Wrap the initialized prompts above as parameters to make them trainable.
35
36        ### Tokenize ###
37        classnames = [name.replace("_", " ") for name in classnames] # 예) "Forest"
38        name_lens = [len(clip.tokenize(name)) for name in classnames]
39        prompts = [prompt_prefix + " " + name + "." for name in classnames] # 예) "A photo of Forest."
40
41        tokenized_prompts = torch.cat([clip.tokenize(p) for p in prompts]) # 예) [49406, 320, 1125, 539...]
42
43
44
45 #####
46 ##### Q1. Fill in the blank #####
47 ##### Define Meta Net #####
48 self.meta_net = nn.Sequential(OrderedDict([
49     ("linear1", nn.Linear(vis_dim, vis_dim // 16)),
50     ("relu", nn.ReLU()),
51     ("linear2", nn.Linear(vis_dim // 16, ctx_dim)),
52     ("sigmoid", nn.Sigmoid())
53 ]))
```

```

50         ("relu", nn.ReLU(inplace=True)),
51         ("linear2", nn.Linear(vis_dim // 16, ctx_dim))
52     ])
53     #####
54     ## Hint: meta network is composed to linear layer, relu activation, and linear layer.
55
56
57
58     if cfg.TRAINER.COOCOOP.PREC == "fp16":
59         self.meta_net.half()
60
61     with torch.no_grad():
62         embedding = clip_model.token_embedding(tokenized_prompts).type(dtype)
63
64     # These token vectors will be saved when in save_model(),
65     # but they should be ignored in load_model() as we want to use
66     # those computed using the current class names
67     self.register_buffer("token_prefix", embedding[:, :1, :]) # SOS
68     self.register_buffer("token_suffix", embedding[:, 1 + n_ctx:, :]) # CLS, EOS
69     self.n_cls = n_cls
70     self.n_ctx = n_ctx
71     self.tokenized_prompts = tokenized_prompts # torch.Tensor
72     self.name_lens = name_lens
73
74     def construct_prompts(self, ctx, prefix, suffix, label=None):
75         # dim0 is either batch_size (during training) or n_cls (during testing)
76         # ctx: context tokens, with shape of (dim0, n_ctx, ctx_dim)
77         # prefix: the sos token, with shape of (n_cls, 1, ctx_dim)
78         # suffix: remaining tokens, with shape of (n_cls, *, ctx_dim)
79
80         if label is not None:
81             prefix = prefix[label]
82             suffix = suffix[label]
83
84         prompts = torch.cat(
85             [
86                 prefix, # (dim0, 1, dim)
87                 ctx, # (dim0, n_ctx, dim)
88                 suffix, # (dim0, *, dim)
89             ],
90             dim=1,
91         )
92
93         return prompts
94
95     def forward(self, im_features):
96         prefix = self.token_prefix
97         suffix = self.token_suffix
98         ctx = self.ctx # (n_ctx, ctx_dim)
99
100
101     #####
102     ##### Q2,3. Fill in the blank #####
103     bias = self.meta_net(im_features) # (batch, ctx_dim)
104     bias = bias.unsqueeze(1) # (batch, 1, ctx_dim)
105     ctx = ctx.unsqueeze(0) # (1, n_ctx, ctx_dim)
106     ctx_shifted = ctx + bias # (batch, n_ctx, ctx_dim)
107     #####
108     #####
109
110
111
112     # Use instance-conditioned context tokens for all classes
113     prompts = []
114     for ctx_shifted_i in ctx_shifted:
115         ctx_i = ctx_shifted_i.unsqueeze(0).expand(self.n_cls, -1, -1)
116         pts_i = self.construct_prompts(ctx_i, prefix, suffix) # (n_cls, n_tkn, ctx_dim)
117         prompts.append(pts_i)
118     prompts = torch.stack(prompts)
119
120
121     return prompts

```

```

1 class CoCoOpCustomCLIP(nn.Module):
2     def __init__(self, cfg, classnames, clip_model):
3         super().__init__()
4         self.prompt_learner = CoCoOpPromptLearner(cfg, classnames, clip_model)
5         self.tokenized_prompts = self.prompt_learner.tokenized_prompts
6         self.image_encoder = clip_model.visual
7         self.text_encoder = TextEncoder(clip_model)
8         self.logit_scale = clip_model.logit_scale
9         self.dtype = clip_model.dtype

```

```

10
11 def forward(self, image, label=None):
12     tokenized_prompts = self.tokenized_prompts
13     logit_scale = self.logit_scale.exp()
14
15     image_features = self.image_encoder(image.type(self.dtype))
16     image_features = image_features / image_features.norm(dim=-1, keepdim=True)
17
18
19     #####
20     ##### Q4. Fill in the blank #####
21     prompts = self.prompt_learner(image_features)
22     #####
23     #####
24
25
26     logits = []
27     for pts_i, imf_i in zip(prompts, image_features):
28         text_features = self.text_encoder(pts_i, tokenized_prompts)
29         text_features = text_features / text_features.norm(dim=-1, keepdim=True)
30         l_i = logit_scale * imf_i @ text_features.t()
31         logits.append(l_i)
32     logits = torch.stack(logits)
33
34     if self.prompt_learner.training:
35         return F.cross_entropy(logits, label)
36
37     return logits

```

❏ Q2. Training CoCoOp

In this task, you will train CoCoOp on the EuroSAT dataset. If your implementation of CoCoOp in Question 1 is correct, the following code should execute without errors. Please submit the execution file so we can evaluate whether your code runs without any issues.

```

1 # Train on the Base Classes Train split and evaluate accuracy on the Base Classes Test split.
2 args.trainer = "CoCoOp"
3 args.train_batch_size = 4
4 args.epoch = 100
5 args.output_dir = "outputs/cocoop"
6
7 args.subsample_classes = "base"
8 args.eval_only = False
9 cocoop_base_acc = main(args)

```

```

🔗 Loading trainer: CoCoOp
Loading dataset: EuroSAT
Reading split from /content/ProMetaR/data/eurosat/split_zhou_EuroSAT.json
Creating a 16-shot dataset
Creating a 4-shot dataset
Saving preprocessed few-shot data to /content/ProMetaR/data/eurosat/split_fewshot/shot_16-seed_1.pkl
SUBSAMPLE BASE CLASSES!
Building transform_train
+ random resized crop (size=(224, 224), scale=(0.08, 1.0))
+ random flip
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
Building transform_test
+ resize the smaller edge to 224
+ 224x224 center crop
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
-----
Dataset      EuroSAT
# classes    5
# train_x    80
# val        20
# test       4,200
-----
Loading CLIP (backbone: ViT-B/16)
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: This DataLoader will create 8 w
  warnings.warn(
Building custom CLIP
Initial context: "a photo of a"
Number of context words (tokens): 4
Turning off gradients in both the image and the text encoder
Parameters to be updated: {'prompt_learner.ctx', 'prompt_learner.meta_net.linear2.weight', 'prompt_learner.meta_net.line
Loading evaluator: Classification
No checkpoint found, train from scratch
/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated
  warnings.warn(
epoch [1/100] batch [20/20] time 0.138 (0.353) data 0.000 (0.032) loss 0.2744 (1.1881) lr 2.5000e-03 eta 0:11:38
epoch [2/100] batch [20/20] time 0.104 (0.129) data 0.000 (0.018) loss 0.8384 (0.8970) lr 2.4994e-03 eta 0:04:11
epoch [3/100] batch [20/20] time 0.095 (0.131) data 0.000 (0.017) loss 0.6382 (0.7859) lr 2.4975e-03 eta 0:04:14

```

```

epoch [4/100] batch [20/20] time 0.095 (0.127) data 0.000 (0.016) loss 0.5044 (0.7151) lr 2.4945e-03 eta 0:04:03
epoch [5/100] batch [20/20] time 0.121 (0.142) data 0.000 (0.015) loss 0.5703 (0.6317) lr 2.4901e-03 eta 0:04:29
epoch [6/100] batch [20/20] time 0.154 (0.192) data 0.000 (0.034) loss 0.6060 (0.6009) lr 2.4846e-03 eta 0:06:01
epoch [7/100] batch [20/20] time 0.097 (0.136) data 0.000 (0.022) loss 0.3853 (0.6638) lr 2.4779e-03 eta 0:04:12
epoch [8/100] batch [20/20] time 0.101 (0.136) data 0.000 (0.019) loss 1.4082 (0.6633) lr 2.4699e-03 eta 0:04:09
epoch [9/100] batch [20/20] time 0.095 (0.131) data 0.000 (0.022) loss 0.1780 (0.4582) lr 2.4607e-03 eta 0:03:58
epoch [10/100] batch [20/20] time 0.141 (0.152) data 0.000 (0.016) loss 1.2285 (0.5051) lr 2.4504e-03 eta 0:04:33
epoch [11/100] batch [20/20] time 0.106 (0.199) data 0.000 (0.036) loss 0.2539 (0.5013) lr 2.4388e-03 eta 0:05:54
epoch [12/100] batch [20/20] time 0.094 (0.126) data 0.000 (0.015) loss 1.1484 (0.4657) lr 2.4261e-03 eta 0:03:41
epoch [13/100] batch [20/20] time 0.096 (0.128) data 0.000 (0.017) loss 0.8467 (0.5009) lr 2.4122e-03 eta 0:03:42
epoch [14/100] batch [20/20] time 0.096 (0.128) data 0.000 (0.023) loss 0.5547 (0.4495) lr 2.3972e-03 eta 0:03:39
epoch [15/100] batch [20/20] time 0.143 (0.152) data 0.000 (0.015) loss 1.0430 (0.5549) lr 2.3810e-03 eta 0:04:18
epoch [16/100] batch [20/20] time 0.136 (0.207) data 0.000 (0.032) loss 1.3906 (0.4799) lr 2.3638e-03 eta 0:05:48
epoch [17/100] batch [20/20] time 0.095 (0.129) data 0.000 (0.017) loss 0.0238 (0.3497) lr 2.3454e-03 eta 0:03:34
epoch [18/100] batch [20/20] time 0.110 (0.130) data 0.000 (0.020) loss 0.1337 (0.2804) lr 2.3259e-03 eta 0:03:33
epoch [19/100] batch [20/20] time 0.094 (0.130) data 0.000 (0.015) loss 1.0420 (0.3864) lr 2.3054e-03 eta 0:03:30
epoch [20/100] batch [20/20] time 0.156 (0.146) data 0.000 (0.024) loss 0.3484 (0.4984) lr 2.2839e-03 eta 0:03:52
epoch [21/100] batch [20/20] time 0.141 (0.193) data 0.000 (0.036) loss 0.8184 (0.3434) lr 2.2613e-03 eta 0:05:05
epoch [22/100] batch [20/20] time 0.096 (0.129) data 0.000 (0.018) loss 0.2090 (0.4361) lr 2.2377e-03 eta 0:03:21

```

```

1 # Accuracy on the New Classes.
2 args.model_dir = "outputs/cocoop"
3 args.output_dir = "outputs/cocoop/new_classes"
4 args.subsample_classes = "new"
5 args.load_epoch = 100
6 args.eval_only = True
7 coop_novel_acc = main(args)

```

```

➔ Loading trainer: CoCoOp
Loading dataset: EuroSAT
Reading split from /content/ProMetaR/data/eurosat/split_zhou_EuroSAT.json
Loading preprocessed few-shot data from /content/ProMetaR/data/eurosat/split_fewshot/shot_16-seed_1.pkl
SUBSAMPLE NEW CLASSES!
Building transform_train
+ random resized crop (size=(224, 224), scale=(0.08, 1.0))
+ random flip
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])
Building transform_test
+ resize the smaller edge to 224
+ 224x224 center crop
+ to torch tensor of range [0, 1]
+ normalization (mean=[0.48145466, 0.4578275, 0.40821073], std=[0.26862954, 0.26130258, 0.27577711])

-----
Dataset      EuroSAT
# classes    5
# train_x    80
# val        20
# test       3,900
-----

Loading CLIP (backbone: ViT-B/16)
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: This DataLoader will create 8 w
arnings.warn(
/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated
warnings.warn(
/content/ProMetaR/dassl/utils/torchtools.py:102: FutureWarning: You are using `torch.load` with `weights_only=False` (th
checkpoint = torch.load(fpath, map_location=map_location)
Building custom CLIP
Initial context: "a photo of a"
Number of context words (tokens): 4
Turning off gradients in both the image and the text encoder
Parameters to be updated: {'prompt_learner.ctx', 'prompt_learner.meta_net.linear2.weight', 'prompt_learner.meta_net.line
Loading evaluator: Classification
Loading weights to prompt_learner from "outputs/cocoop/prompt_learner/model.pth.tar-100" (epoch = 100)
Evaluate on the *test* set
100%|██████████| 39/39 [01:01<00:00, 1.57s/it]=> result
* total: 3,900
* correct: 1,687
* accuracy: 43.3%
* error: 56.7%
* macro_f1: 39.0%

```

✓ Q3. Analyzing the results of CoCoOp

Compare the results of CoCoOp with those of CoOp that we trained in Lab Session 4. Discuss possible reasons for the performance differences observed between CoCoOp and CoOp.

The results show that CoCoOp performs well on the base classes with an accuracy of 90.8%, which is strong. However, on new classes, its performance drops dramatically to 43.3%. Compared to CoOp, CoOp has lower accuracy on new classes. This demonstrates CoOp's fixed prompt tuning generalizes better to unseen categories in this case.

Possible reasons for the performance difference could be that CoCoOp relies on dynamic, image-conditioned prompts and this causes overfitting to visual features of base classes. Thus, its ability to handle unseen categories. In contrast CoOp has a better generalization across

class distributions.